

Matplotlib_and_seabon

October 16, 2021

```
[ ]: import numpy as np
import pandas as pd
```

```
[ ]: import seaborn as sns
```

```
[ ]: import matplotlib.pyplot as plt
```

```
[ ]: stocks= pd.read_csv('stocks.csv')
```

```
[ ]: stocks.head()
```

```
[ ]:
```

	Date	AAPL	IBM	CSCO	MSFT
0	2000-01-03	111.937502	116.0000	108.0625	116.5625
1	2000-01-04	102.500003	112.0625	102.0000	112.6250
2	2000-01-05	103.999997	116.0000	101.6875	113.8125
3	2000-01-06	94.999998	114.0000	100.0000	110.0000
4	2000-01-07	99.500001	113.5000	105.8750	111.4375


```
[ ]: stocks['Date']=pd.to_datetime(stocks['Date'])
stocks.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3521 entries, 0 to 3520
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Date    3521 non-null     datetime64[ns]
1   AAPL    3521 non-null     float64
2   IBM     3521 non-null     float64
3   CSCO    3521 non-null     float64
4   MSFT    3521 non-null     float64
dtypes: datetime64[ns](1), float64(4)
memory usage: 137.7 KB
```


0.1 Line Plot

```
[ ]: sns.lineplot(x='Date',y='AAPL',data=stocks)
```

```
[ ]: <AxesSubplot:xlabel='Date', ylabel='AAPL'>
```



```
[ ]: # Now using matplotliblib
plt.plot(stocks['Date'],stocks['AAPL'])
plt.title('apple stock')
plt.show()
```



```
[ ]: plt.subplot(2,2,1)
plt.plot(stocks['Date'],stocks['AAPL'])
plt.title('apple stock')


plt.subplot(2,2,2)
plt.plot(stocks['Date'],stocks['MSFT'])
plt.title('Microsoft stock')

plt.subplot(2,2,3)
plt.plot(stocks['Date'],stocks['CSCO'])
plt.title('Cisco stock')

plt.subplot(2,2,4)
plt.plot(stocks['Date'],stocks['IBM'])
plt.title('IBM stock')


plt.tight_layout()
plt.show
```

```
[ ]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[ ]: stocks= stocks.set_index('Date')
stocks.plot()
```

```
[ ]: <AxesSubplot:xlabel='Date'>
```



0.2 Scatterplots

captures relationship between 2 continues variables

```
[ ]: cars=pd.read_csv('cars.csv')
cars.head()
```

```
[ ]:
```

	name	sports_car	suv	wagon	minivan	pickup	\
0	Chevrolet Aveo 4dr	False	False	False	False	False	
1	Chevrolet Aveo LS 4dr hatch	False	False	False	False	False	
2	Chevrolet Cavalier 2dr	False	False	False	False	False	
3	Chevrolet Cavalier 4dr	False	False	False	False	False	
4	Chevrolet Cavalier LS 2dr	False	False	False	False	False	

	all_wheel	rear_wheel	Price	Dealer_Cost	Engine_size	cylenders	\
0	False	False	11690	10965	1.6	4	
1	False	False	12585	11802	1.6	4	
2	False	False	14610	13697	2.2	4	
3	False	False	14810	13884	2.2	4	
4	False	False	16385	15357	2.2	4	

	horsepower	city_miles_per_galloon	highway_miles_per_Gallon	weight	\
0	103	28.0	34.0	2370.0	
1	103	28.0	34.0	2348.0	
2	140	26.0	37.0	2617.0	
3	140	26.0	37.0	2676.0	
4	140	26.0	37.0	2617.0	

	base_wheeeel	length	width
0	98.0	167.0	66.0
1	98.0	153.0	66.0
2	104.0	183.0	69.0
3	104.0	183.0	68.0
4	104.0	183.0	69.0

```
[ ]: cars.shape
```

```
[ ]: (428, 19)
```

```
[ ]: cylenders= cars.cylenders.value_counts()
cylenders
```

```
[ ]: 6      190
      4      136
      8       87
      5        7
      12       3
     -1        2
      10       2
      3        1
```

Name: cylenders, dtype: int64

notice that cylinders 4, 6 and 8 is the meajority

```
[ ]: cars_common= cars[cars.cylenders.isin([4,6,8])]
cars_common
```

```
[ ]:
      name  sports_car  suv  wagon  minivan  \
0      Chevrolet Aveo 4dr      False  False  False  False
1      Chevrolet Aveo LS 4dr hatch      False  False  False  False
2      Chevrolet Cavalier 2dr      False  False  False  False
3      Chevrolet Cavalier 4dr      False  False  False  False
4      Chevrolet Cavalier LS 2dr      False  False  False  False
..      ...
423      Nissan Titan King Cab XE      False  False  False  False
424      Subaru Baja      False  False  False  False
425      Toyota Tacoma      False  False  False  False
426      Toyota Tundra Regular Cab V6      False  False  False  False
427      Toyota Tundra Access Cab V6 SR5      False  False  False  False
```

	pickup	all_wheel	rear_wheel	Price	Dealer_Cost	Engine_size \
0	False	False	False	11690	10965	1.6
1	False	False	False	12585	11802	1.6
2	False	False	False	14610	13697	2.2
3	False	False	False	14810	13884	2.2
4	False	False	False	16385	15357	2.2
..
423	True	True	False	26650	24926	5.6
424	True	True	False	24520	22304	2.5
425	True	False	True	12800	11879	2.4
426	True	False	True	16495	14978	3.4
427	True	True	False	25935	23520	3.4


	cylinders	horsepower	city_miles_per_galloon	highway_miles_per_Gallon \
0	4	103	28.0	34.0
1	4	103	28.0	34.0
2	4	140	26.0	37.0
3	4	140	26.0	37.0
4	4	140	26.0	37.0
..
423	8	305	14.0	18.0
424	4	165	21.0	28.0
425	4	142	22.0	27.0
426	6	190	16.0	20.0
427	6	190	14.0	17.0

	weight	base_wheeeel	length	width
0	2370.0	98.0	167.0	66.0
1	2348.0	98.0	153.0	66.0
2	2617.0	104.0	183.0	69.0
3	2676.0	104.0	183.0	68.0
4	2617.0	104.0	183.0	69.0
..
423	5287.0	140.0	NaN	NaN
424	3485.0	104.0	NaN	NaN
425	2750.0	103.0	NaN	NaN
426	3925.0	128.0	NaN	NaN
427	4435.0	128.0	NaN	NaN

[413 rows x 19 columns]


```
[ ]: sns.scatterplot(x='horsepower',y='Price',data=cars_common)
```

```
[ ]: <AxesSubplot:xlabel='horsepower', ylabel='Price'>
```




```
[ ]: sns.scatterplot(x='horsepower',y='city_miles_per_galloon',data=cars_common)
```

```
[ ]: <AxesSubplot:xlabel='horsepower', ylabel='city_miles_per_galloon'>
```



```
[ ]: sns.  
      ↳scatterplot(x='horsepower',y='city_miles_per_galloon',data=cars_common,hue='cylinders')
```

```
[ ]: <AxesSubplot:xlabel='horsepower', ylabel='city_miles_per_galloon'>
```



```
[ ]: cars_common[['cylinders']].info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 413 entries, 0 to 427  
Data columns (total 1 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   cylinders    413 non-null    int64  
dtypes: int64(1)  
memory usage: 22.6 KB
```

Now change to categories, because cylinders is integer

```
[ ]: cars_common['cylinders']=cars_common['cylinders'].astype('category')  
cars_common[['cylinders']].info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 413 entries, 0 to 427  
Data columns (total 1 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   cylinders    413 non-null    category
```



```

---  -----  -----
0    cylenders  413 non-null    category
dtypes: category(1)
memory usage: 19.9 KB

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.


Cylenders is now categorical type

```

[ ]: sns.
      ↳scatterplot(x='horsepower',y='city_miles_per_galloon',data=cars_common,hue='cylenders')

[ ]: <AxesSubplot:xlabel='horsepower', ylabel='city_miles_per_galloon'>

```



0.3 Count plot

```
[ ]: ##### barplots
retail= pd.read_csv('online_retail2.csv')
retail= retail.drop_duplicates()
retail= retail.dropna(axis=0,how='any')
retail.Country.value_counts()
```

```
[ ]: United Kingdom      716115
      Germany            17339
      EIRE               16014
      France            13897
      Netherlands       5137
      Spain             3754
      Belgium           3110
      Switzerland       3058
      Portugal          2414
      Australia         1890
      Channel Islands   1646
      Italy             1507
      Sweden            1343
      Norway            1308
      Cyprus            1157
      Finland           1049
      Austria           938
      Denmark           797
      Greece            663
      Japan             565
      USA               535
      Poland            527
      Unspecified       521
      United Arab Emirates 386
      Singapore         346
      Israel            321
      Malta             299
      Canada            228
      Iceland           222
      Lithuania         154
      RSA               123
      Brazil            94
      Thailand          76
      Korea             63
      European Community 61
      Bahrain           59
      West Indies       54
      Lebanon           45
      Nigeria           30
```

```
Czech Republic      30
Saudi Arabia        10
Name: Country, dtype: int64
```

```
[ ]: retail.shape
```

```
[ ]: (797885, 8)
```

```
[ ]: ire_Fran= retail[retail.Country.isin(['EIRE','France'])]
ire_Fran
```

```
[ ]:      Invoice StockCode      Description  Quantity \
71      489439      22065      CHRISTMAS PUDDING TRINKET POT      12
72      489439      22138      BAKING SET 9 PIECE RETROSPOT      9
73      489439      22139      RETRO SPOT TEA SET CERAMIC 11 PC      9
74      489439      22352      LUNCHBOX WITH CUTLERY RETROSPOT      12
75      489439      85014A      BLACK/BLUE DOTS RUFFLED UMBRELLA      3
...
1067366  581587      22899      CHILDREN'S APRON DOLLY GIRL      6
1067367  581587      23254      CHILDRENS CUTLERY DOLLY GIRL      4
1067368  581587      23255      CHILDRENS CUTLERY CIRCUS PARADE      4
1067369  581587      22138      BAKING SET 9 PIECE RETROSPOT      3
1067370  581587      POST      POSTAGE      1
```

```
      InvoiceDate  Price  Customer ID  Country
71      2009-12-01 09:28:00    1.45      12682.0  France
72      2009-12-01 09:28:00    4.95      12682.0  France
73      2009-12-01 09:28:00    4.95      12682.0  France
74      2009-12-01 09:28:00    2.55      12682.0  France
75      2009-12-01 09:28:00    5.95      12682.0  France
...
1067366  2011-12-09 12:50:00    2.10      12680.0  France
1067367  2011-12-09 12:50:00    4.15      12680.0  France
1067368  2011-12-09 12:50:00    4.15      12680.0  France
1067369  2011-12-09 12:50:00    4.95      12680.0  France
1067370  2011-12-09 12:50:00   18.00      12680.0  France
```

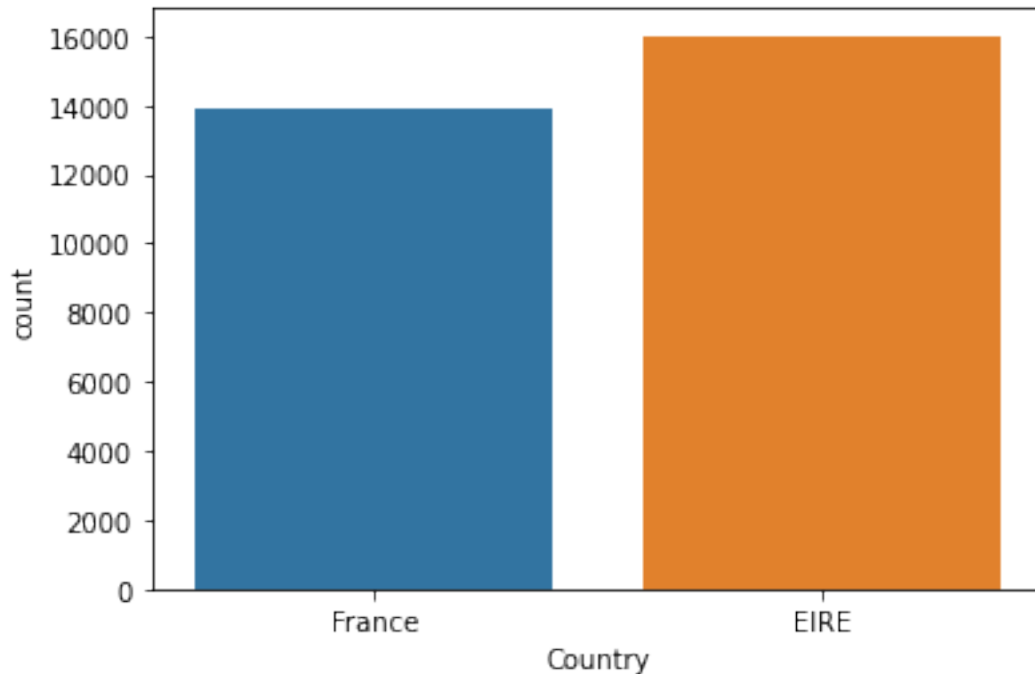
```
[29911 rows x 8 columns]
```

```
[ ]: ire_Fran.shape
```

```
[ ]: (29911, 8)
```

```
[ ]: sns.countplot(x='Country',data= ire_Fran)
```

```
[ ]: <AxesSubplot:xlabel='Country', ylabel='count'>
```



```
[ ]: ire_Fran['InvoiceDate'].head()
```

```
[ ]: 71    2009-12-01 09:28:00
      72    2009-12-01 09:28:00
      73    2009-12-01 09:28:00
      74    2009-12-01 09:28:00
      75    2009-12-01 09:28:00
      Name: InvoiceDate, dtype: object
```

0.4 Barplot

```
[ ]: from datetime import datetime
```

```
[ ]: ire_Fran[['InvoiceDate']].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29911 entries, 71 to 1067370
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  -
0   InvoiceDate  29911 non-null  object
dtypes: object(1)
memory usage: 1.5+ MB
```

This is an object, convert it to datetime

```
[ ]: ire_Fran['InvoiceDate']=pd.to_datetime(ire_Fran['InvoiceDate'])
```

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
"""Entry point for launching an IPython kernel.

```
[ ]: ire_Fran[['InvoiceDate']].info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 29911 entries, 71 to 1067370  
Data columns (total 1 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   InvoiceDate  29911 non-null  datetime64[ns]  
dtypes: datetime64[ns](1)  
memory usage: 1.5 MB
```

Now it is of type datetime

```
[ ]: ire_Fran['dayofweek']= ire_Fran['InvoiceDate'].dt.dayofweek  
ire_Fran['dayofweek']
```

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
"""Entry point for launching an IPython kernel.

```
[ ]: 71      1  
     72      1  
     73      1  
     74      1  
     75      1  
     ..  
    1067366  4  
    1067367  4  
    1067368  4  
    1067369  4  
    1067370  4  
Name: dayofweek, Length: 29911, dtype: int64
```

Create a dictionary and label days of week

```
[ ]: week_dict= { 0 : 'Monday', 1: 'Tuesday', 2: 'Wednesday', 3: 'thursday', 4: 'Friday', 5: 'Saturday', 6: 'Sunday' }
week_dict
```

```
[ ]: {0: 'Monday',
      1: 'Tuesday',
      2: 'Wednesday',
      3: 'thursday',
      4: 'Friday',
      5: 'Saturday',
      6: 'Sunday'}
```

```
[ ]: ire_Fran['day_of_wk']=ire_Fran['dayofweek'].map(week_dict)
ire_Fran['day_of_wk'].head()
```

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 """Entry point for launching an IPython kernel.

```
[ ]: 71    Tuesday
      72    Tuesday
      73    Tuesday
      74    Tuesday
      75    Tuesday
      Name: day_of_wk, dtype: object
```

```
[ ]: grouped= ire_Fran.groupby(['Country','day_of_wk']).agg(total_quantity=np.sum).reset_index()
grouped=grouped[grouped.total_quantity >=0]
```


```
[ ]: grouped
```

```
[ ]:   Country  day_of_wk  total_quantity
0    EIRE    Friday         45835
1    EIRE    Monday         47061
2    EIRE    Sunday         31796
3    EIRE    Tuesday         54930
4    EIRE  Wednesday         57443
5    EIRE  thursday         72652
6  France    Friday         34814
7  France    Monday        111764
8  France    Sunday         15321
10 France  Wednesday         35035
```

11 France thursday 41462


```
[ ]: sns.barplot(x='day_of_wk',y='total_quantity',data=grouped)
```

```
[ ]: <AxesSubplot:xlabel='day_of_wk', ylabel='total_quantity'>
```



```
[ ]: sns.barplot(x='day_of_wk',y='total_quantity',data=grouped,hue='Country')
```

```
[ ]: <AxesSubplot:xlabel='day_of_wk', ylabel='total_quantity'>
```



```
[ ]: ire_Fran.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29911 entries, 71 to 1067370
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Invoice          29911 non-null  object
1   StockCode       29911 non-null  object
2   Description     29911 non-null  object
3   Quantity        29911 non-null  int64
4   InvoiceDate     29911 non-null  datetime64[ns]
5   Price           29911 non-null  float64
6   Customer ID    29911 non-null  float64
7   Country         29911 non-null  object
8   dayofweek       29911 non-null  int64
9   day_of_wk       29911 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(2), object(5)
memory usage: 3.5+ MB
```



0.5 Distribution Plot

```
[ ]: iris= pd.read_csv('iris.csv')
iris.head()
```

```
[ ]:      sepal_length  sepal_width  petal_length  petal_width  species
0          5.1           3.5           1.4           0.2   setosa
1          4.9           3.0           1.4           0.2   setosa
2          4.7           3.2           1.3           0.2   setosa
3          4.6           3.1           1.5           0.2   setosa
4          5.0           3.6           1.4           0.2   setosa
```

```
[ ]: sns.kdeplot(iris.sepal_length)
```

```
[ ]: <AxesSubplot:>
```



Sepal length pair each type of flower

```
[ ]: iris.species.value_counts()
```


```
[ ]: versicolor    50
     virginica     50
     setosa        50
     Name: species, dtype: int64
```

```
[ ]: iris.sepal_length[iris.species=='virginica'].count()
```

```
[ ]: 50
```

```
[ ]: sns.kdeplot(iris.sepal_length[iris.species=='virginica'],label='virginica')  
sns.kdeplot(iris.sepal_length[iris.species=='setosa'],label='setosa')  
sns.kdeplot(iris.sepal_length[iris.species=='versicolor'],label='versicolor')
```

```
[ ]: <AxesSubplot:>
```




in general, The sepal length of the setosa is lower than the sepal length of versicolor and sepal length of versicolor is lower virginica

```
[ ]: iris.sepal_length.describe()
```

```
[ ]: count    150.000000  
mean        5.843333  
std         0.828066  
min         4.300000  
25%         5.100000  
50%         5.800000  
75%         6.400000  
max         7.900000  
Name: sepal_length, dtype: float64
```

```
[ ]: sns.kdeplot(iris.sepal_width[iris.species=='virginica'],label='Verginica')  
sns.kdeplot(iris.sepal_width[iris.species=='setosa'],label='setosa')  
sns.kdeplot(iris.sepal_width[iris.species=='versicolor'],label='versicolor')
```

```
[ ]: <AxesSubplot:>
```




The setosa is more higher in width than the other 2.

0.6 Boxplot

we will see the outliers and percentiles using the Boxplot


```
[ ]: sns.boxplot(x ="species", y="sepal_length",data=iris)
```

```
[ ]: <AxesSubplot:xlabel='species', ylabel='sepal_length'>
```



```
[ ]: sns.boxplot(x= 'species',y='sepal_width',data=iris)
```

```
[ ]: <AxesSubplot:xlabel='species', ylabel='sepal_width'>
```



```
[ ]: cars_common
```

```
[ ]:
      name sports_car  suv  wagon  minivan  \
0      Chevrolet Aveo 4dr      False  False  False  False
1      Chevrolet Aveo LS 4dr hatch      False  False  False  False
2      Chevrolet Cavalier 2dr      False  False  False  False
3      Chevrolet Cavalier 4dr      False  False  False  False
4      Chevrolet Cavalier LS 2dr      False  False  False  False
..      ...
423      Nissan Titan King Cab XE      False  False  False  False
424      Subaru Baja      False  False  False  False
425      Toyota Tacoma      False  False  False  False
426      Toyota Tundra Regular Cab V6      False  False  False  False
427      Toyota Tundra Access Cab V6 SR5      False  False  False  False

      pickup  all_wheel  rear_wheel  Price  Dealer_Cost  Engine_size  cylenders  \
0      False      False      False  11690      10965      1.6      4
1      False      False      False  12585      11802      1.6      4
2      False      False      False  14610      13697      2.2      4
3      False      False      False  14810      13884      2.2      4
4      False      False      False  16385      15357      2.2      4
..      ...
423      True      True      False  26650      24926      5.6      8
424      True      True      False  24520      22304      2.5      4
425      True      False      True  12800      11879      2.4      4
426      True      False      True  16495      14978      3.4      6
427      True      True      False  25935      23520      3.4      6

      horsepower  city_miles_per_galloon  highway_miles_per_Gallon  weight  \
0      103      28.0      34.0  2370.0
1      103      28.0      34.0  2348.0
2      140      26.0      37.0  2617.0
3      140      26.0      37.0  2676.0
4      140      26.0      37.0  2617.0
..      ...
423      305      14.0      18.0  5287.0
424      165      21.0      28.0  3485.0
425      142      22.0      27.0  2750.0
426      190      16.0      20.0  3925.0
427      190      14.0      17.0  4435.0

      base_wheeeel  length  width
0      98.0      167.0      66.0
1      98.0      153.0      66.0
2      104.0      183.0      69.0
```

```


3          104.0    183.0    68.0
4          104.0    183.0    69.0
..          ...      ...      ...
423        140.0     NaN     NaN
424        104.0     NaN     NaN
425        103.0     NaN     NaN
426        128.0     NaN     NaN
427        128.0     NaN     NaN

```

[413 rows x 19 columns]

```
[ ]: sns.boxplot(x='cylinders', y='horsepower', data=cars_common)
```

```
[ ]: <AxesSubplot:xlabel='cylinders', ylabel='horsepower'>
```




0.7 Histograms

In Histogram, x is always continues and we are doing frequency count of the continues variable.

```
[ ]: plt.hist(x= iris.sepal_length, bins=30)
```


```
[ ]: (array([ 4.,  1.,  4.,  2.,  5., 16.,  9.,  4.,  1.,  6., 13.,  8.,  7.,
            3.,  6., 10.,  9.,  7.,  5.,  2., 11.,  4.,  1.,  1.,  4.,  1.,
            0.,  1.,  4.,  1.]),
      array([4.3 , 4.42, 4.54, 4.66, 4.78, 4.9 , 5.02, 5.14, 5.26, 5.38, 5.5 ,
```

```
5.62, 5.74, 5.86, 5.98, 6.1 , 6.22, 6.34, 6.46, 6.58, 6.7 , 6.82,
6.94, 7.06, 7.18, 7.3 , 7.42, 7.54, 7.66, 7.78, 7.9 ]),
<BarContainer object of 30 artists>)
```




```
[ ]: sns.distplot(iris.sepal_length)
```

```
[ ]: <AxesSubplot:xlabel='sepal_length'>
```



```
[ ]: sns.distplot(iris.sepal_length,bins=30,kde=False)
```

```
[ ]: <AxesSubplot:xlabel='sepal_length'>
```



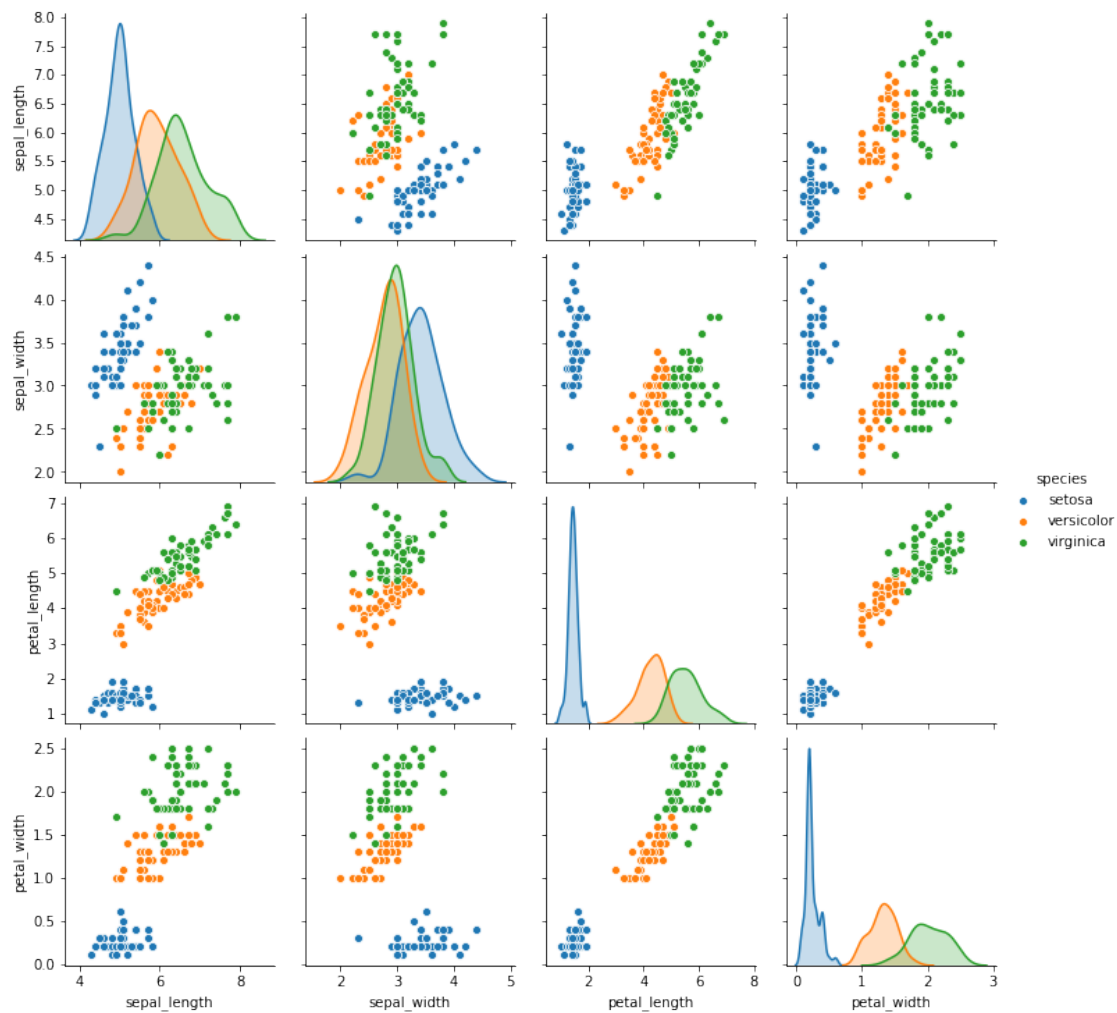
0.8 Pairplots

It does multiple comparisons between many variables as long as we specify which variables.

E.g Compare Setosa, virginica and versicula

```
[ ]: sns.pairplot(iris,hue='species')
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x1de93994bc8>
```



As seen petal lengths of setosa is very low, if a flower has petal length of 2, most definitely, it could be a setosa.

if a flower has petal width from 0 to 1, most definitely it might be a setosa.

petal widths more than 2.0, most definitely it is a virginica.

```
[ ]: cars_common.columns
```

```
[ ]: Index(['name', 'sports_car', 'suv', 'wagon', 'minivan', 'pickup', 'all_wheel',  
          'rear_wheel', 'Price', 'Dealer_Cost', 'Engine_size', 'cylenders',  
          'horsepower', 'city_miles_per_galloon', 'highway_miles_per_Gallon',  
          'weight', 'base_wheeeel', 'length', 'width'],  
          dtype='object')
```

```
[ ]: cars_common.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 413 entries, 0 to 427  
Data columns (total 19 columns):  
#   Column                                Non-Null Count  Dtype    
---  ---                                -  
0   name                                413 non-null   object   
1   sports_car                         413 non-null   bool     
2   suv                                413 non-null   bool     
3   wagon                              413 non-null   bool     
4   minivan                           413 non-null   bool     
5   pickup                             413 non-null   bool     
6   all_wheel                          413 non-null   bool     
7   rear_wheel                         413 non-null   bool     
8   Price                              413 non-null   int64    
9   Dealer_Cost                        413 non-null   int64    
10  Engine_size                        413 non-null   float64  
11  cylenders                          413 non-null   category  
12  horsepower                         413 non-null   int64    
13  city_miles_per_galloon             402 non-null   float64  
14  highway_miles_per_Gallon           402 non-null   float64  
15  weight                             411 non-null   float64  
16  base_wheeeel                      411 non-null   float64  
17  length                             387 non-null   float64  
18  width                              387 non-null   float64  
dtypes: bool(7), category(1), float64(7), int64(3), object(1)  
memory usage: 58.2+ KB
```

```
[ ]: cars_common.iloc[0:2]
```

```
[ ]:      name  sports_car  suv  wagon  minivan  pickup  \  
0      Chevrolet Aveo 4dr      False  False  False      False  False  
1  Chevrolet Aveo LS 4dr hatch      False  False  False      False  False  
  
      all_wheel  rear_wheel  Price  Dealer_Cost  Engine_size  cylenders  \  
0      False      False  11690      10965      1.6      4  
1      False      False  12585      11802      1.6      4  
  
      horsepower  city_miles_per_galloon  highway_miles_per_Gallon  weight  \  

```

0	103	28.0	34.0	2370.0
1	103	28.0	34.0	2348.0

	base_wheeeel	length	width
0	98.0	167.0	66.0
1	98.0	153.0	66.0

```
[ ]: sns.
      ↪ pairplot(cars_common[['horsepower', 'sports_car', 'Price', 'city_miles_per_galloon', 'cylenders'
```

```

      ↪ -----
      ↪
      ↪ ValueError                                Traceback (most recent call
      ↪ last)

```

```

      ↪ ~\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py in
      ↪ kdensityfft(X, kernel, bw, weights, gridsize, adjust, clip, cut, retgrid)
      ↪ 450     try:
      ↪ --> 451         bw = float(bw)
      ↪ 452     except:

```

ValueError: could not convert string to float: 'scott'

During handling of the above exception, another exception occurred:

```

      ↪ RuntimeError                                Traceback (most recent call
      ↪ last)

```

```

      ↪ <ipython-input-69-11e2a3f1cfe8> in <module>
      ↪ ----> 1 sns.
      ↪ ↪ pairplot(cars_common[['horsepower', 'sports_car', 'Price', 'city_miles_per_galloon', 'cylenders'
```

```

      ↪ ~\Anaconda3\lib\site-packages\seaborn\axisgrid.py in pairplot(data, hue,
      ↪ hue_order, palette, vars, x_vars, y_vars, kind, diag_kind, markers, height,
      ↪ aspect, corner, dropna, plot_kws, diag_kws, grid_kws, size)
      ↪ 2119         diag_kws.setdefault("shade", True)
      ↪ 2120         diag_kws["legend"] = False
      ↪ -> 2121         grid.map_diag(kdeplot, **diag_kws)
      ↪ 2122
      ↪ 2123     # Maybe plot on the off-diagonals

```

```

~\Anaconda3\lib\site-packages\seaborn\axisgrid.py in map_diag(self,
↳func, **kwargs)
1488         data_k = utils.remove_na(data_k)
1489
-> 1490         func(data_k, label=label_k, color=color, **kwargs)
1491
1492         self._clean_axis(ax)

~\Anaconda3\lib\site-packages\seaborn\distributions.py in kdeplot(data,
↳data2, shade, vertical, kernel, bw, gridsize, cut, clip, legend, cumulative,
↳shade_lowest, cbar, cbar_ax, cbar_kws, ax, **kwargs)
703         ax = _univariate_kdeplot(data, shade, vertical, kernel, bw,
704                                 gridsize, cut, clip, legend, ax,
--> 705                                 cumulative=cumulative, **kwargs)
706
707         return ax

~\Anaconda3\lib\site-packages\seaborn\distributions.py in
↳_univariate_kdeplot(data, shade, vertical, kernel, bw, gridsize, cut, clip,
↳legend, ax, cumulative, **kwargs)
293         x, y = _statsmodels_univariate_kde(data, kernel, bw,
294                                             gridsize, cut, clip,
--> 295                                             cumulative=cumulative)
296     else:
297         # Fall back to scipy if missing statsmodels

~\Anaconda3\lib\site-packages\seaborn\distributions.py in
↳_statsmodels_univariate_kde(data, kernel, bw, gridsize, cut, clip, cumulative)
365         fft = kernel == "gau"
366         kde = smnp.KDEUnivariate(data)
--> 367         kde.fit(kernel, bw, fft, gridsize=gridsize, cut=cut, clip=clip)
368         if cumulative:
369             grid, y = kde.support, kde.cdf

~\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py in
↳fit(self, kernel, bw, fft, weights, gridsize, adjust, cut, clip)
138         density, grid, bw = kdensityfft(endog, kernel=kernel,
↳bw=bw,
139                                     adjust=adjust, weights=weights,
↳gridsize=gridsize,
--> 140                                     clip=clip, cut=cut)
141     else:

```

```


142             density, grid, bw = kdensity(endog, kernel=kernel, bw=bw,

~\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py in
↪kdensityfft(X, kernel, bw, weights, gridsize, adjust, clip, cut, retgrid)
    451         bw = float(bw)
    452     except:
--> 453         bw = bandwidths.select_bandwidth(X, bw, kern) # will
↪cross-val fit this pattern?
    454         bw *= adjust
    455

~\Anaconda3\lib\site-packages\statsmodels\nonparametric\bandwidths.py in
↪select_bandwidth(x, bw, kernel)
    172         # eventually this can fall back on another selection
↪criterion.
    173         err = "Selected KDE bandwidth is 0. Cannot estimate density."
--> 174         raise RuntimeError(err)
    175     else:
    176         return bandwidth

```

RuntimeError: Selected KDE bandwidth is 0. Cannot estimate density.



[]:

[]: